

HP-Bot

A retrieval-augmented chatbot that answers **only** from the Harry Potter dataset.

What the project asked for

You gave us a small Harry Potter dataset — 20 question/answer pairs and 130 raw passages — and ten requirements. They group into four jobs:

01–06

Behave

Six graded rules: refuse off-topic and out-of-data questions, greet without leaking internals, resist jailbreaks, remember the conversation, ignore format demands.

07–08

Retrieve

A two-stage FAISS design — a question cache that skips the LLM entirely, then a full-data index that feeds the most similar chunks to the model.

09

Interface

A simple way to talk to the bot. The brief said it can be minimal, so a Streamlit chat with a panel that shows how each answer was reached.

10

Report

A short write-up: the tech stack, a flow diagram, and the parts that were difficult and the parts that were enjoyable.


Technology stack

Everything is pip-installable and runs on a laptop CPU. No GPU, no Docker, no build step.

Language	Python 3.11+	Fits the FAISS + sentence-transformers ecosystem.
Interface	Streamlit	A chat UI in one command — no HTML or CSS.
Embeddings	sentence-transformers · all-MiniLM-L6-v2	Smallest model that still works well — about 80 MB, CPU-only.
Vector index	FAISS (CPU)	Course-mandated. Cosine via inner product on normalised vectors.
Sparse retrieval	rank-bm25	Blended with dense search so rare names still surface.
LLM	OpenRouter — glm-4.5-air (free) by default	One API, any model. Seven-model fallback chain behind it.
Config	python-dotenv (.env)	API key and thresholds in one file. Model is swappable.
Tests	YAML cases + Playwright	40 adversarial cases, plus a script that drives the live UI.

Flow diagram: how a message is answered

Most messages never reach the LLM — the guard and the question cache answer them first.



The guard and the cache resolve the easy cases — instantly and for free.

Only genuinely new questions reach the model — then the answer is appended to memory.

Two FAISS indices, two jobs

The smart-chunking part of the brief — and the cheapest path to a correct answer.

INDEX A

Questions only

Indexed on the dataset's questions alone.

Embed the new question, take the top-1 cosine match.

Score ≥ 0.85 → return the stored answer verbatim.

No API call. Instant, deterministic, free, and it can't hallucinate — a cached answer is fixed text.

INDEX B

All of the data

Indexed on everything — questions, answers, passages.

Used only when Index A misses (score < 0.85).

Hybrid score: $0.7 \cdot \text{dense FAISS} + 0.3 \cdot \text{BM25}$.

Top-5 chunks become the LLM's context. BM25 catches rare proper nouns dense embeddings sometimes miss.

Net effect: common questions are answered instantly and for free; only genuinely new questions pay for an LLM call.

The six rules — solved in the prompt

src/prompts.py is the most important file. Every graded rule is a numbered instruction; a regex guard adds defence in depth.

1

Off-topic → refuse

Not about the dataset? Reply the exact string “I cannot answer that..”

2

Not in data → refuse

Harry Potter but not in our corpus? Same refusal — no parametric fallback.

3

Greet, don't leak

A whitelist answers “hi” / “who are you?” but never reveals how it works.

4

Resist jailbreaks

“Ignore previous instructions”, fake-admin, “write me code” — refused twice over.

5

Remember the turn


Last 5 turns kept, so “how old is he?” resolves against the prior question.

6

Ignore format demands

“Answer in ten words / in French / as JSON” — the demand is dropped.

The interface



Why Streamlit


The brief said the UI can be minimal, so I picked something I could ship in one command. A chat widget and a sidebar, no HTML or CSS.

The detail I like


Every reply carries an expandable “**retrieval details**” panel: was the answer caught by the guard, served from the cache, or written by the LLM? You can see exactly how each answer was reached — useful for debugging and for proving the bot does what I claim.

It greets, and it answers for free

A greeting handled by the prompt's whitelist, then a question served straight from the cache.




- **Rule 3** "hi" gets a friendly whitelisted reply — no refusal.




- **Rules 7 & 8** "Buckbeak?" hits Index A — source: cache, no LLM.

Two ways to say no

An off-topic question, and a classic injection attack. Same refusal, different paths.




- **Rule 1** off-topic → the exact two-dot refusal string.




- **Rule 4** injection caught by the regex guard — source: guard.

It remembers, and it won't be reformatted

The pronoun “she” resolves to the previous turn; a “reply in French” demand is ignored.



- **Rule 5** two turns — “she” resolves to Hermione.



- **Rule 6** “reply in French” ignored — answers in English.

40 adversarial cases — 40 pass

40

of 40 pass

0 regressions

0 mismatches

0 errors

RULE	WHAT IT TESTS	PASS
1	Out-of-scope refusal — capitals, code, recipes, LOTR, Marvel	8 / 8
2	Out-of-knowledge refusal — HP facts not in the corpus	6 / 6
3	Greeting / identity / capability whitelist	6 / 6
4	Jailbreak / injection / internals disclosure	10 / 10
5	Pronoun resolution across multi-turn follow-ups	5 / 5
6	Format / style manipulation — 10 words, JSON, French, pirate	5 / 5

Reproduce: `python -m tests.run_eval` · a second runner retries each case 3× to absorb free-tier provider noise.

What was hard, what I enjoyed

DIFFICULT

The exact refusal string. Two dots — not one, not three, not a full sentence. Fixed by quoting it verbatim and telling the model to copy it character-for-character.

Greet vs. don't-reveal. “Who are you?” must answer; “how do you work?” must refuse. A whitelist plus a forbidden-topics list separates them.

Flaky free-tier LLMs. Errors, null content, ignored temperature. A seven-model fallback chain keeps the bot answering and surfaces real outages honestly.

ENJOYED

Writing the eval suite. Reading prompt-injection write-ups, encoding them as labelled cases, and watching the pass rate climb from 22 to 40 as the prompt tightened.

The two-stage retrieval. Most tutorials call the LLM every time. A question cache makes common answers instant, free — and impossible to hallucinate.

The cold-start solve. Streamlit's `cache_resource` turned a 30-second start into an instant chat after the first run. Small change, big feel.

THANK YOU

Happy to take questions —
architecture, prompt, eval, anything.

GRADING CHECKLIST

SUBMISSION.md

FULL WRITE-UP

REPORT.md · HP-Bot-Report.pdf

PER-CASE EVAL

REPORT-eval-new-corpus.md

SLIDE DECK

HP-Bot-presentation.pdf